

# EdgeRDV: A Framework for Edge Workload Management at Scale

Gloire Rubambiza (gloire@cs.cornell.edu)

IEEE Edge - 07/04/2023 - Chicago, IL

Co-authors: Braulio Dumba (IBM Research), Andrew J. Anderson (IBM Research),  
Hakim Weatherspoon (Cornell University)

Motivation



State-of-the-art

Methodology

Analysis/Evaluation



# Motivation

- Edge Computing
  - Supports manufacturing, smart homes, supply chains, etc
  - Reduces latency and bandwidth consumption
  - Novel applications such as AI inference at the edge



# Edge Computing Challenges

- How do we manage application life-cycle at scale?
- Can we do so with limited cloud connectivity?
- How do we sustain operation when connectivity is lost?
- How do we gracefully recover from endpoint failures?



# Edge Computing Challenges

- How do we manage application life-cycle at scale?
- Can we do so with limited cloud connectivity?
- How do we sustain operation when connectivity is lost?
- How do we gracefully recover from endpoint failures?

**Research Question:** how do we efficiently manage the application lifecycle on hundreds of thousands of endpoints?

Motivation

State-of-the-art

Methodology


Analysis/Evaluation

# Edge Workload Management in Resource-Constrained Environments

- Scaling Community Networks (CNs) with Community Cellular Manager. Hasan et al. NSDI. 2019.
- Energy-efficient Computing for Wildlife Tracking: Design Tradeoffs and Early Experiences with ZebraNet. Juang et al. ASPLOS. 2002.
- Visage: Enabling Timely Analytics for Drone Imagery. Jha et al. MobiCom. 2021.
- The Akamai Network: A Platform for High-Performance Internet Applications. Nygren et al. UMass Amherst. 2010.
- Experience in Implementing a Non-IP Routing Protocol VIRO in GENI. Dumba et al. 2014.

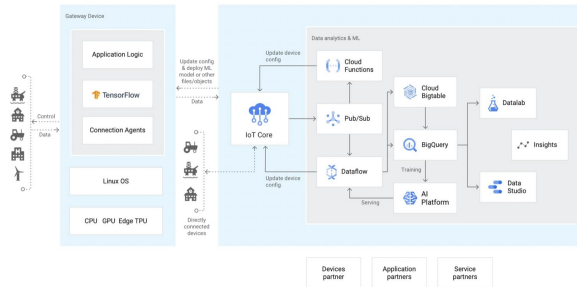
# Industrial IoT Management Platforms

## Azure IoT Edge




- Use cases
  - **Manufacturing**
  - **Electric**
  - **Automation**

## Google IoT Core



- Use cases
  - **Transportation**
  - **Utilities**
  - **Healthcare**


## AWS IoT Core



- Use cases
  - **Manufacturing**
  - **Smart homes**
  - **Supply chain**




# Industrial IoT Management Platforms

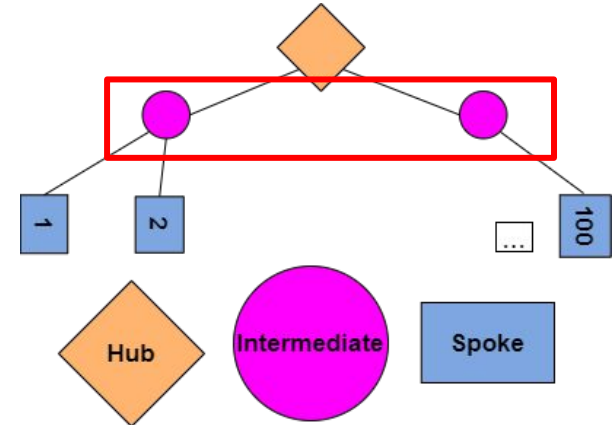


Google IoT, AWS IoT, IBM EAM

# Industrial IoT Management Platforms




Google IoT, AWS IoT, IBM EAM




Azure IoT

# Industrial IoT Management Platforms




# Industrial IoT Management Platforms

- Application life-cycle management at scale
- Limited connectivity
- Disconnected operation
- Failure recovery



# Industrial IoT Management Platforms

- Application life-cycle management at scale
  - **Bootstrapping IoT runtimes** 👍
  - **Layered deployments** 👍
- Limited connectivity
  - **Communication modalities** 👍
- Disconnected operation
  - **Offline cache of messages** 👍
- Failure recovery



# Industrial IoT Management Platforms

- Application life-cycle management at scale

- Bootstrapping IoT runtimes 👍
- Layered deployments 👍
- **Min. bandwidth consumption** ❌

- Limited connectivity


- Communication modalities 👍
- **Peer-to-peer communication** ❌

- Disconnected operation

- Offline cache of messages 👍

- Failure recovery

- **Caching workloads** ❌
- **Avoid single point of failure** ❌



Motivation


State-of-the-art

**Methodology**

**Analysis/Evaluation**

# Proposed Solution – Key Idea

## Scaling Edge Deployments using Rendezvous Nodes

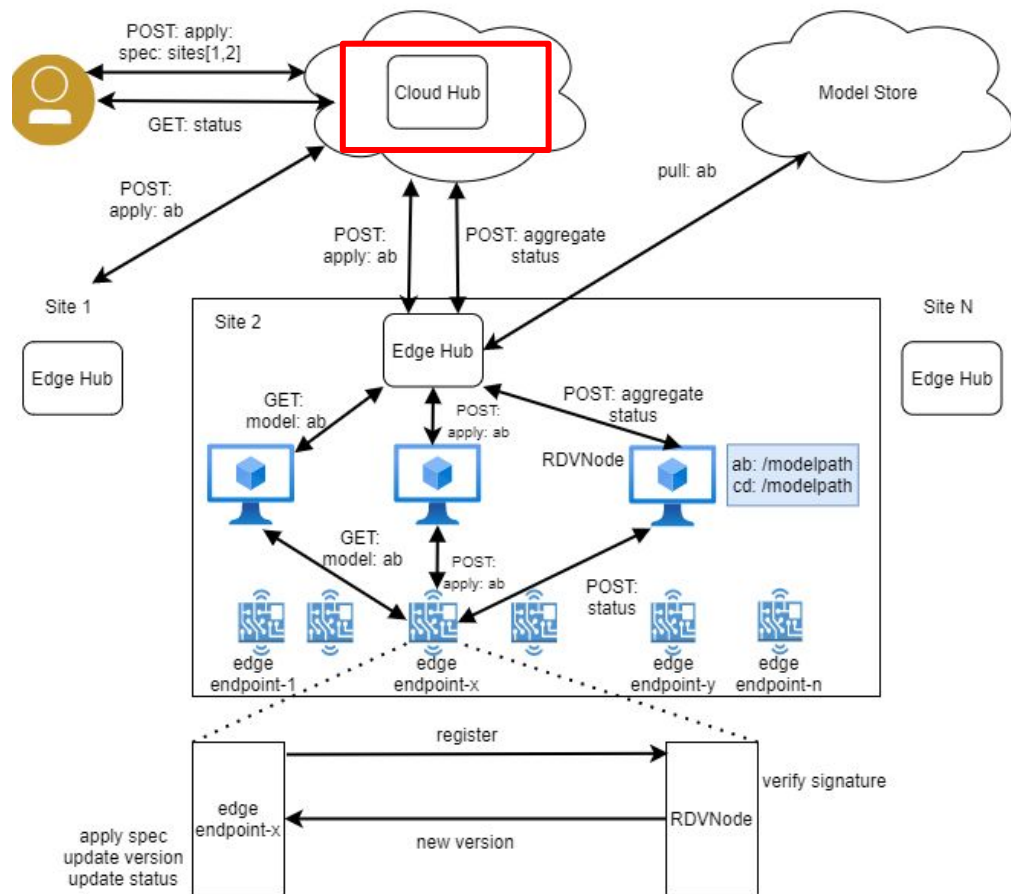


VIRO: A Scalable, Robust and Namespace Independent **V**irtual **I**d **R**outing for Future Networks, Jain *et al.*, IEEE INFOCOMM, 2011

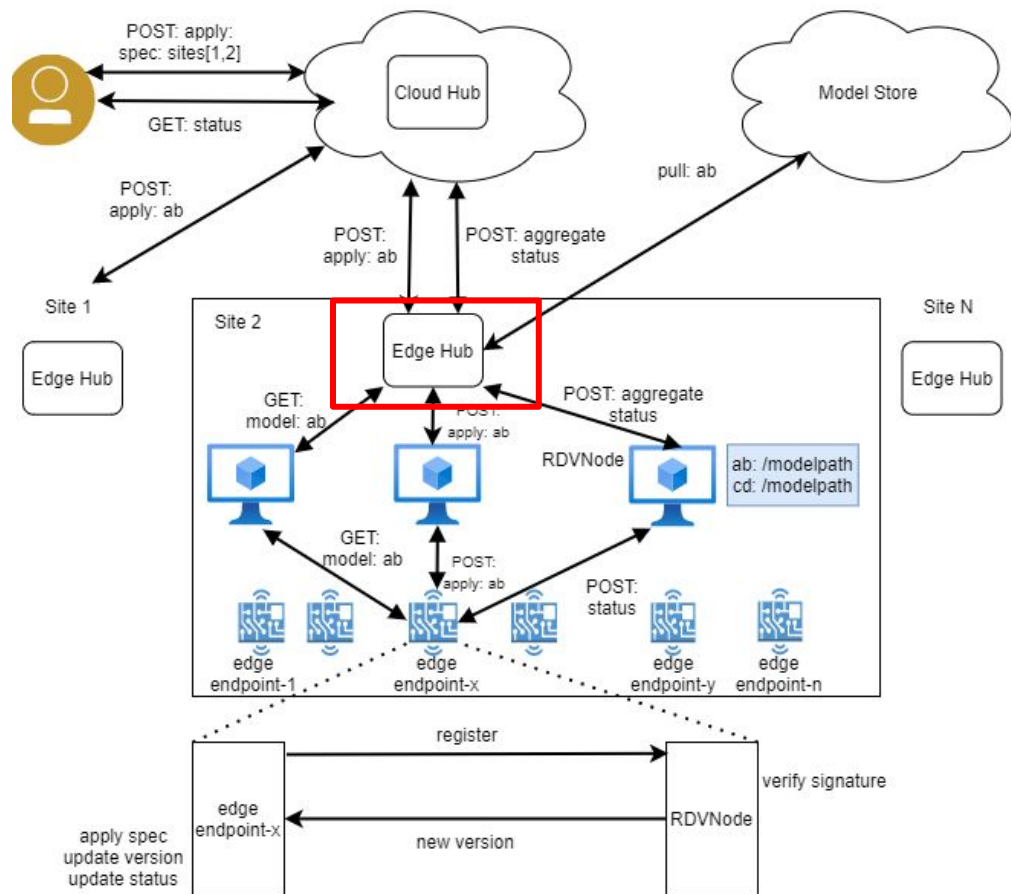
Experience in Implementing a Non-IP Routing Protocol VIRO in GENI, Dumba *et al.*, IEEE ICNP, 2014



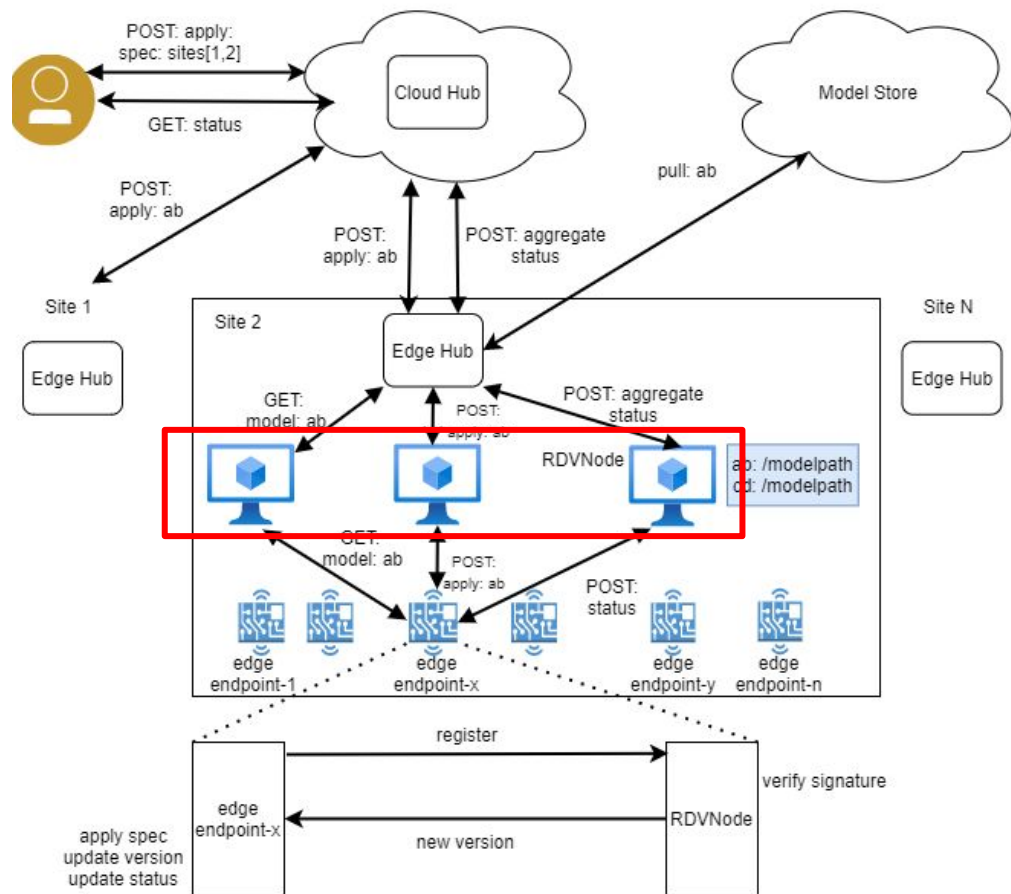
# EdgeRDV Benefits: Scalability



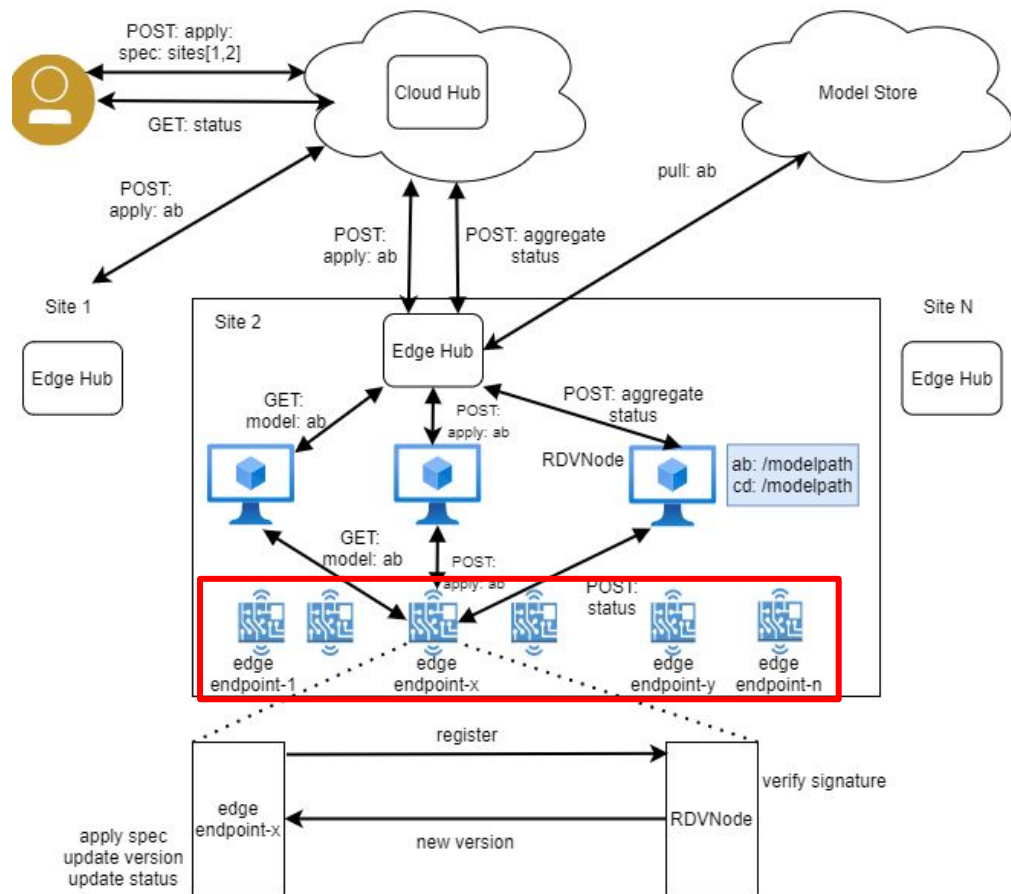
# EdgeRDV Benefits: Scalability



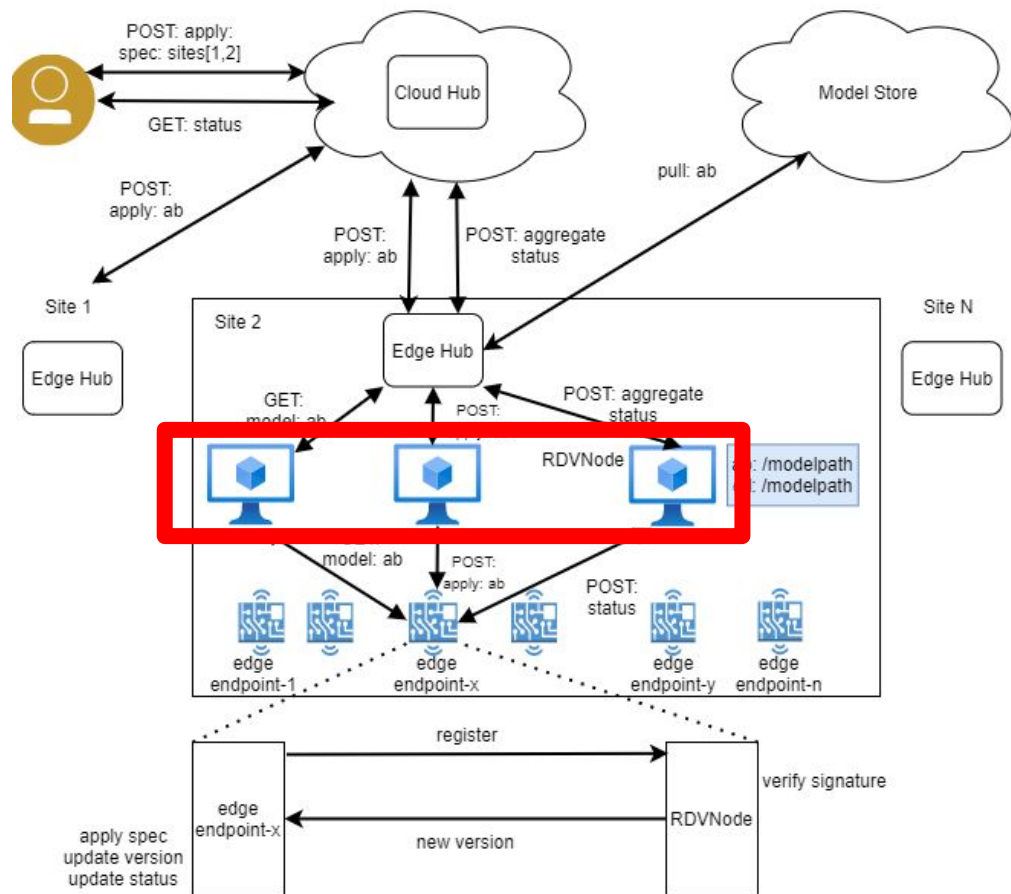
# EdgeRDV Benefits: Scalability



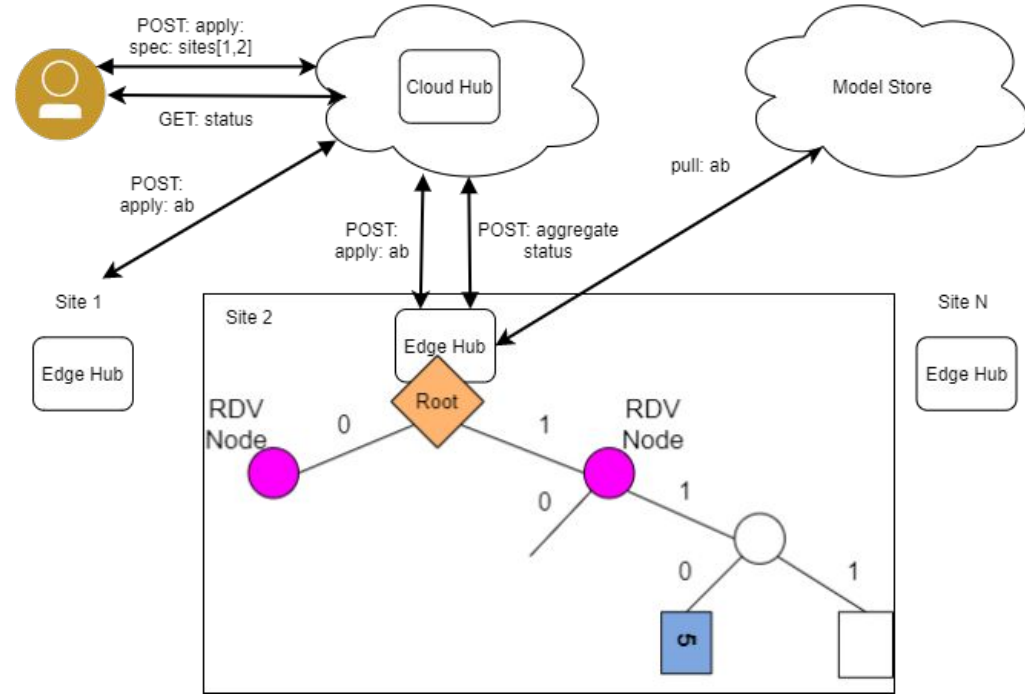
# EdgeRDV Benefits: Scalability



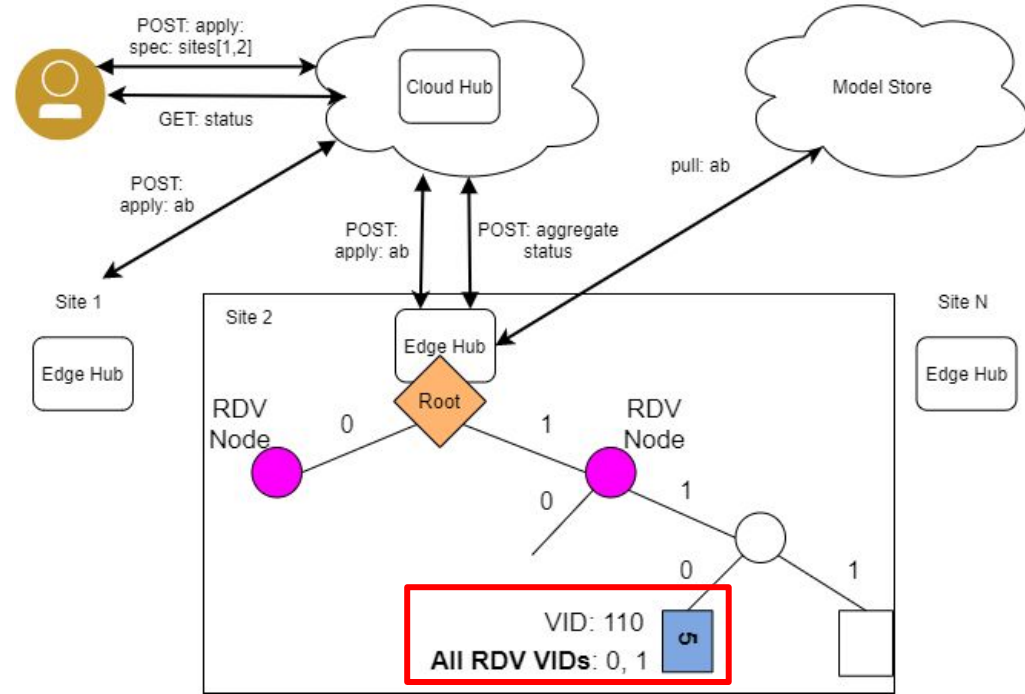
# EdgeRDV Benefits: Scalability



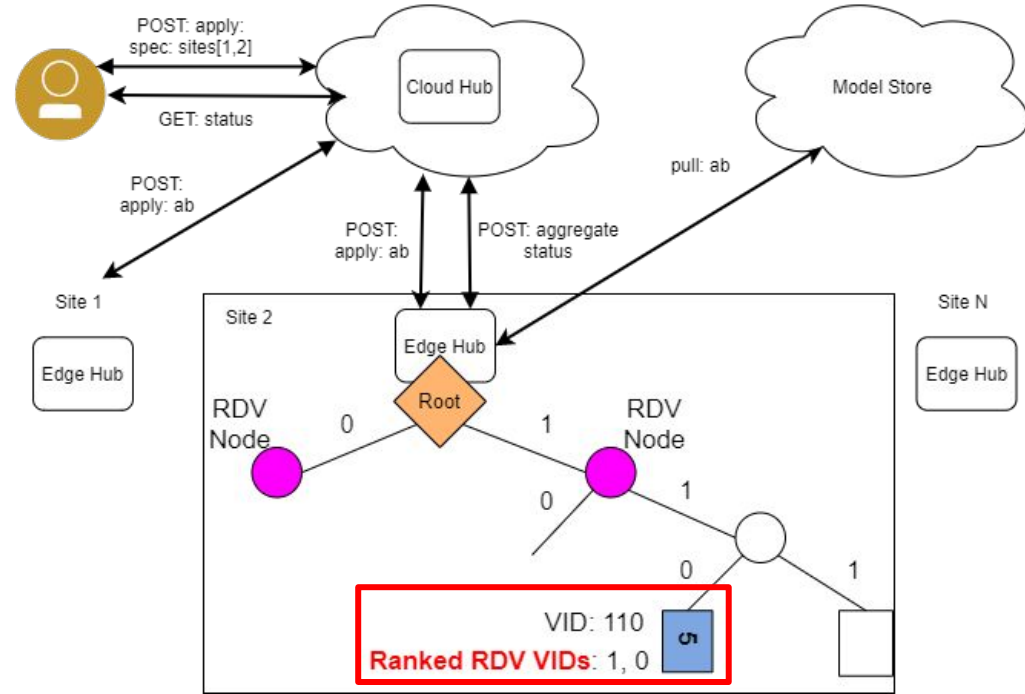
# EdgeRDV Benefits: Scalability



# EdgeRDV Benefits: Scalability



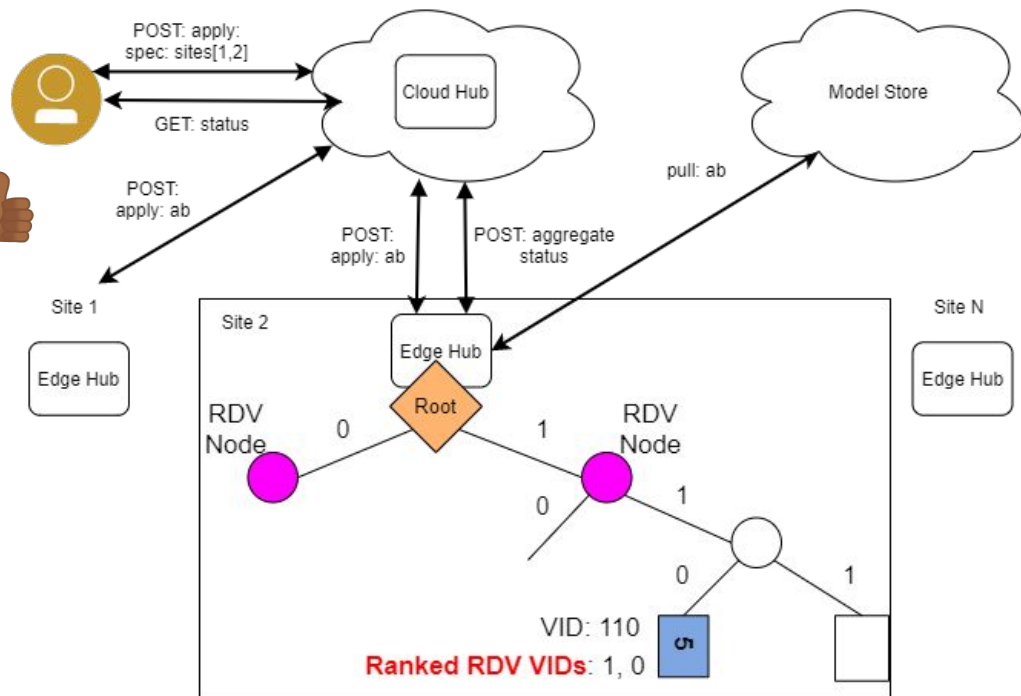
# EdgeRDV Benefits: Scalability





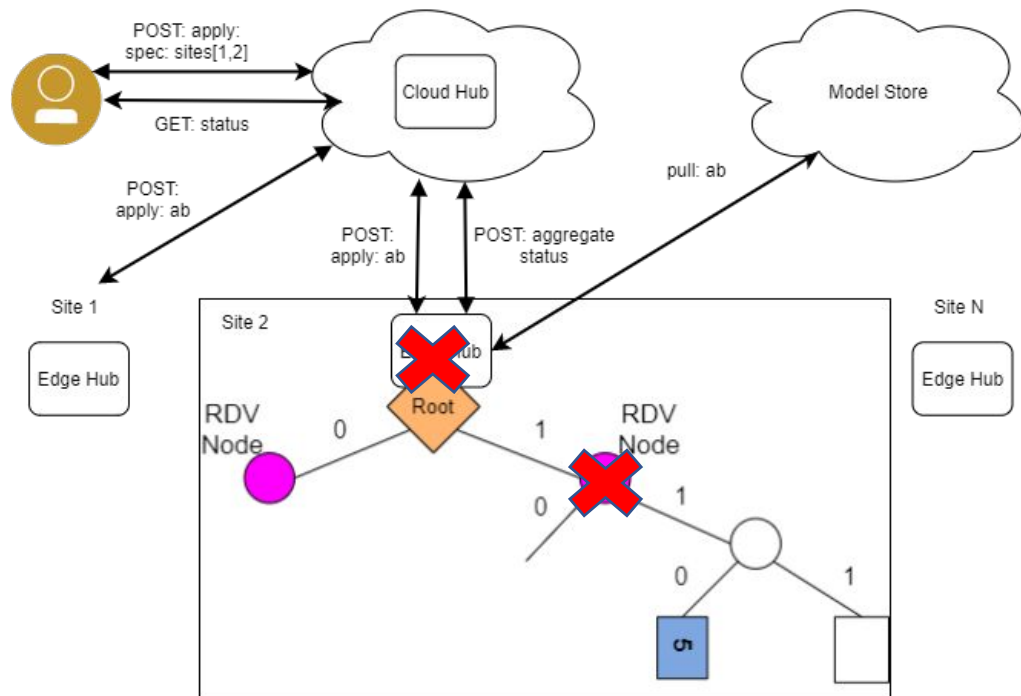
# EdgeRDV Benefits: Scalability

- Priority list of RDV nodes 👍
- Minimal edge hub overhead 👍

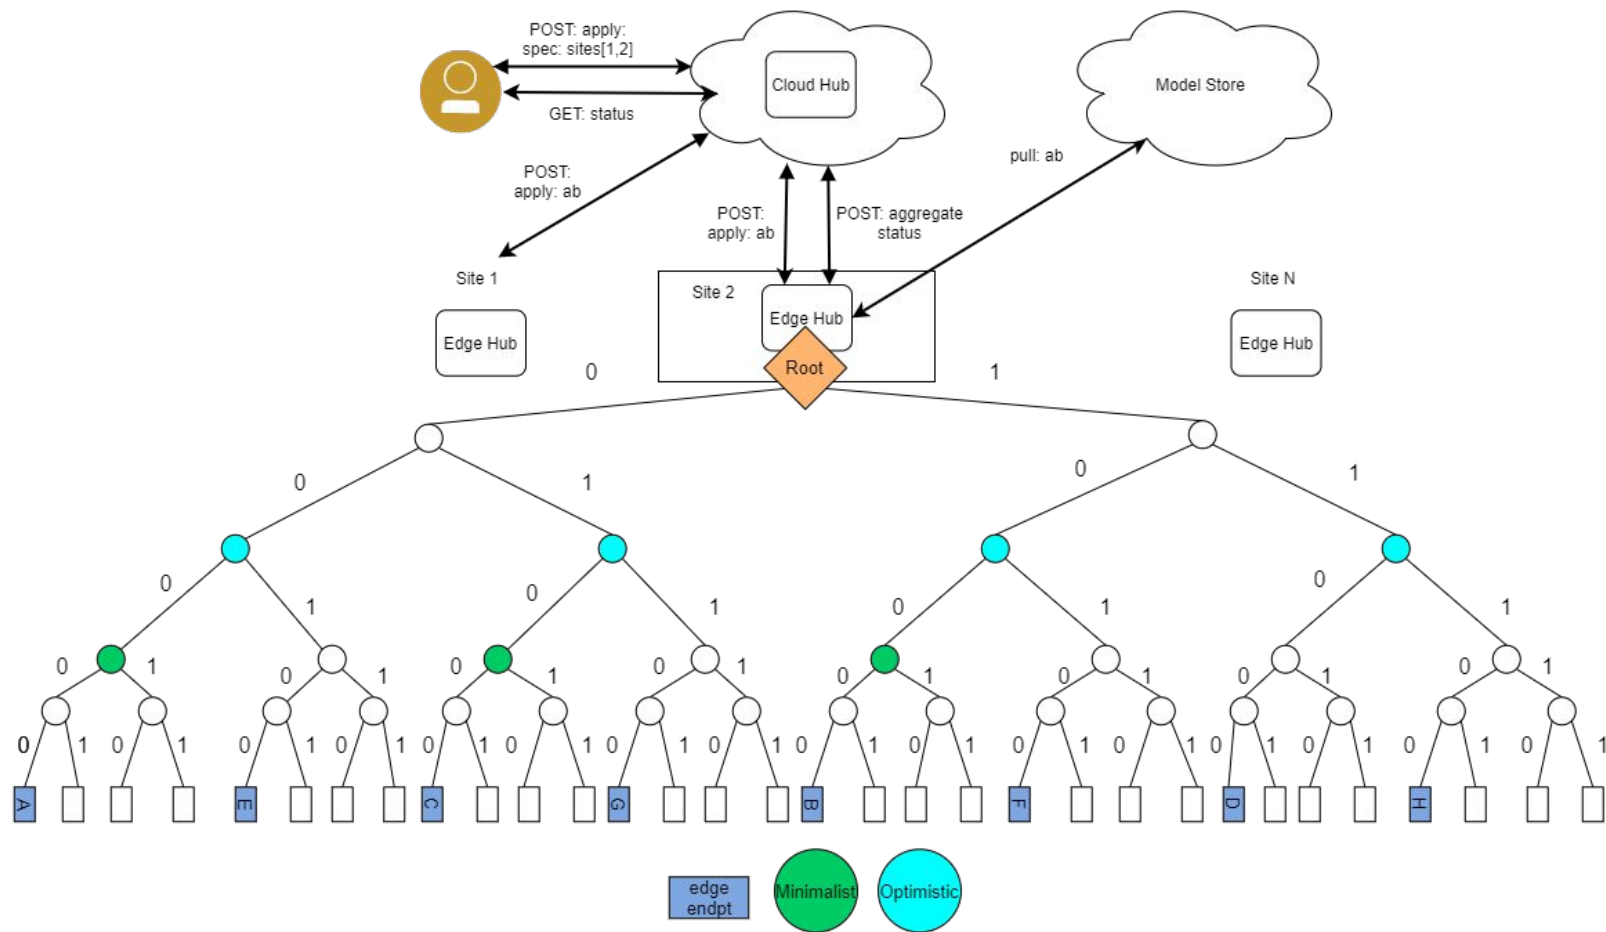


# EdgeRDV Benefits: Multi-level Failure Resilience

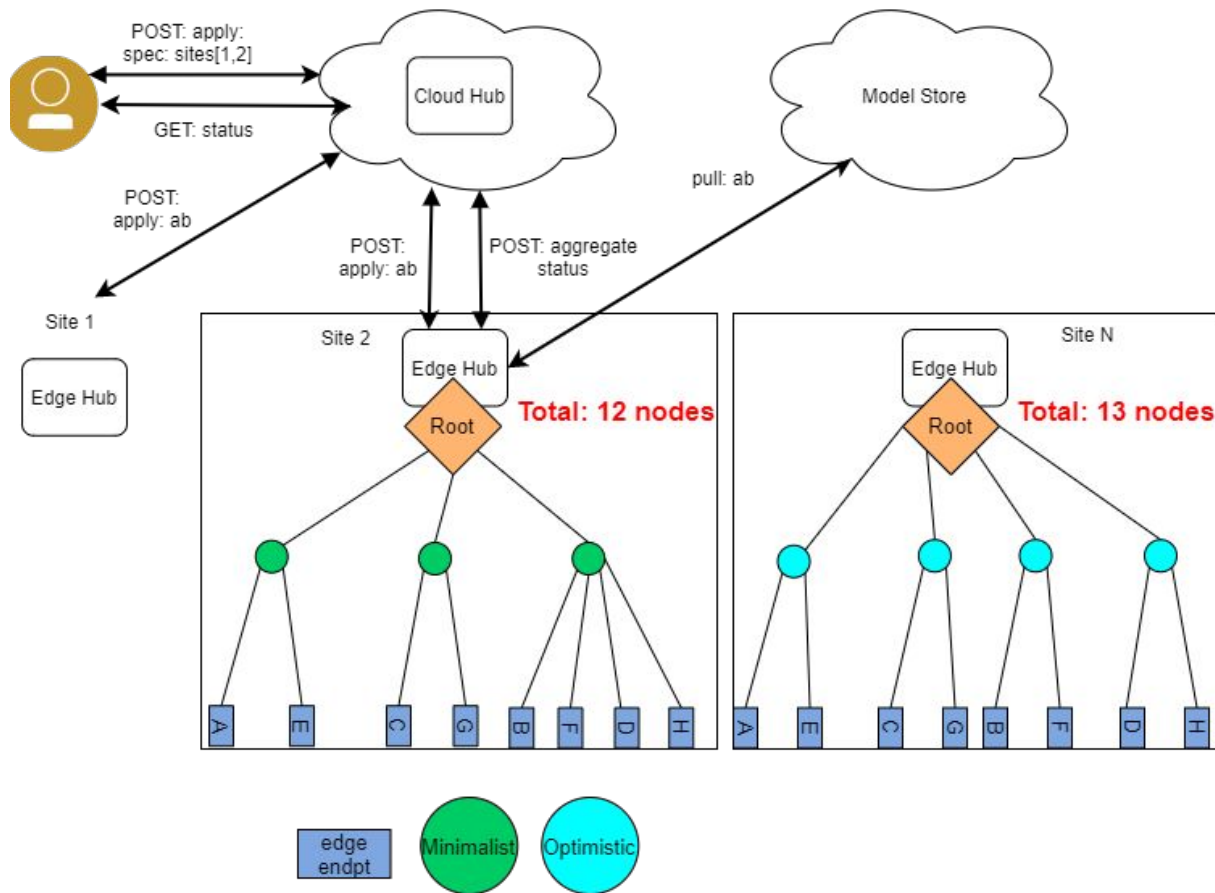
- Multi-level caching 👍
- Multiple RDV nodes 👍
- Ranked list of backup RDVs 👍



# Challenge with RDV nodes - How many and where?



# Challenge with RDV nodes - How many and where?




# Simulation Results

## Keeping the network balanced

256 Edge Endpoints, 5 RDVs

Edge Endpt-to-RDV Assignments as a Function of Tree Depth




**Minimalist**

# Simulation Results

## Keeping the network balanced

256 Edge Endpoints, 5 RDVs


Edge Endpt-to-RDV Assignments as a Function of Tree Depth



**Minimalist**

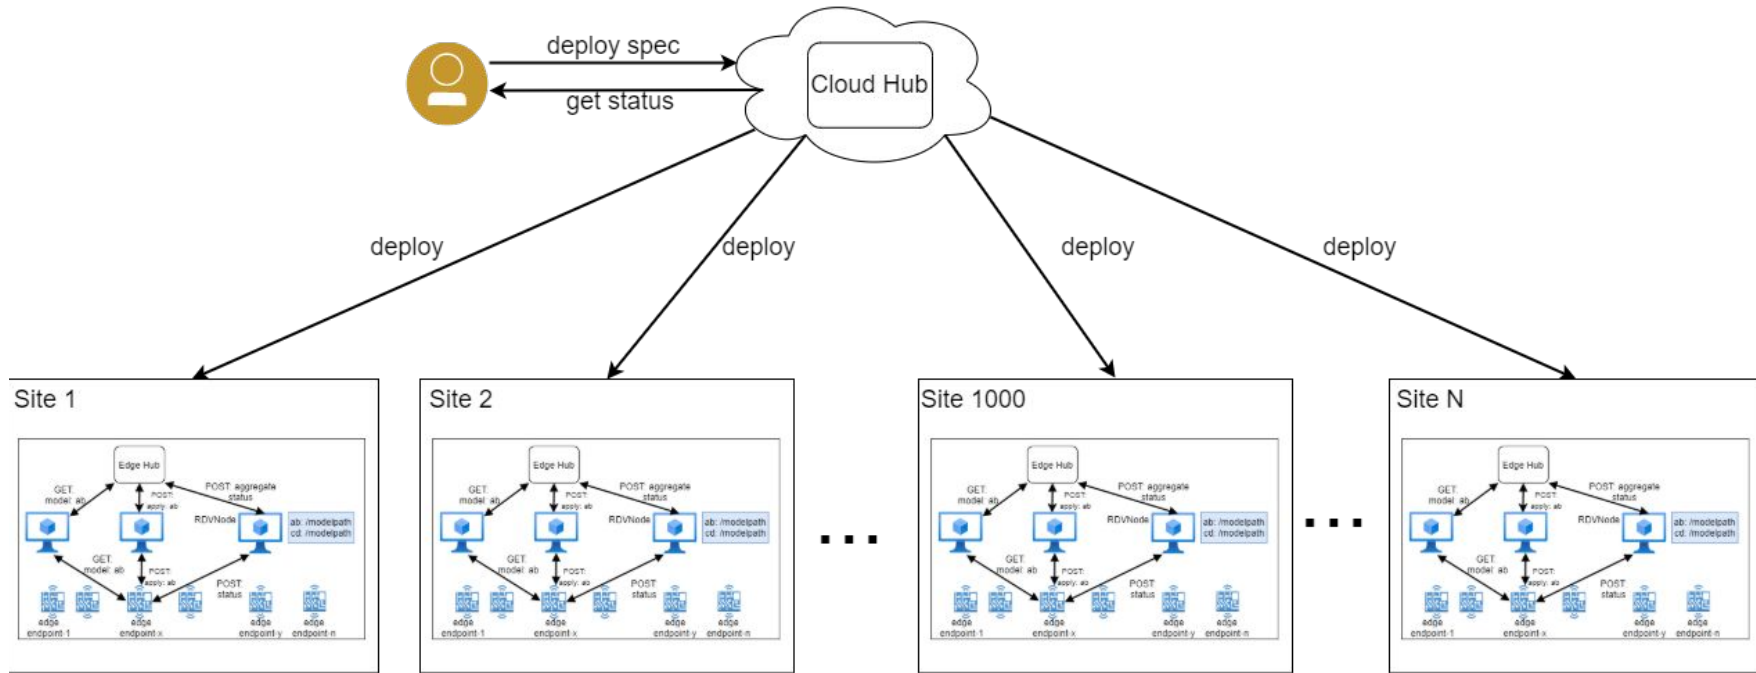
256 Edge Endpoints, 4 RDVs

Edge Endpt-to-RDV Assignments as a Function of RDV Depth




**Optimistic**

# EdgeRDV at Scale



# EdgeRDV at Scale





Motivation

State-of-the-art

Methodology

**Analysis/Evaluation**

# Experiments

- **Setup**

- Docker Linux container (Ubuntu 20.04)
- RDV node coverage: 10%
- Run components as processes on the same host
- Scalability analysis up to 667K endpoints

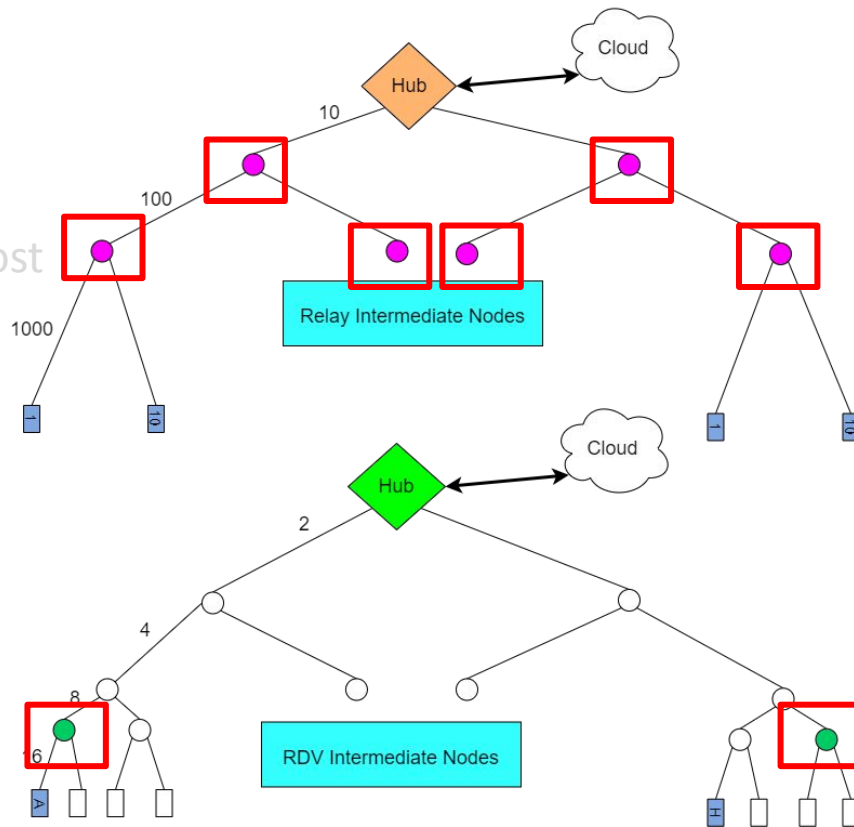
# Experiments

## • Setup

- Docker Linux container (Ubuntu 20.04)
- RDV node coverage: 10%
- Run components as processes on the same host
- Scalability analysis up to 667K endpoints

## • Metrics

- Intermediate nodes



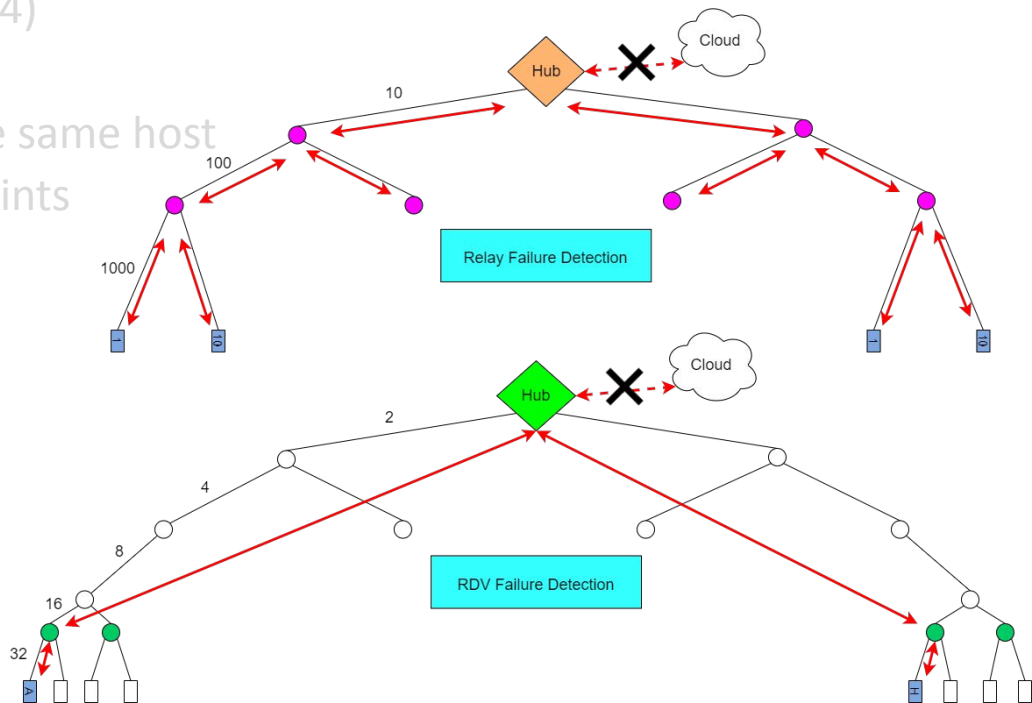
# Experiments

## • Setup

- Docker Linux container (Ubuntu 20.04)
- RDV node coverage: 10%
- Run components as processes on the same host
- Scalability analysis up to 667K endpoints

## • Metrics

- Intermediate nodes
- **Total physical edges**



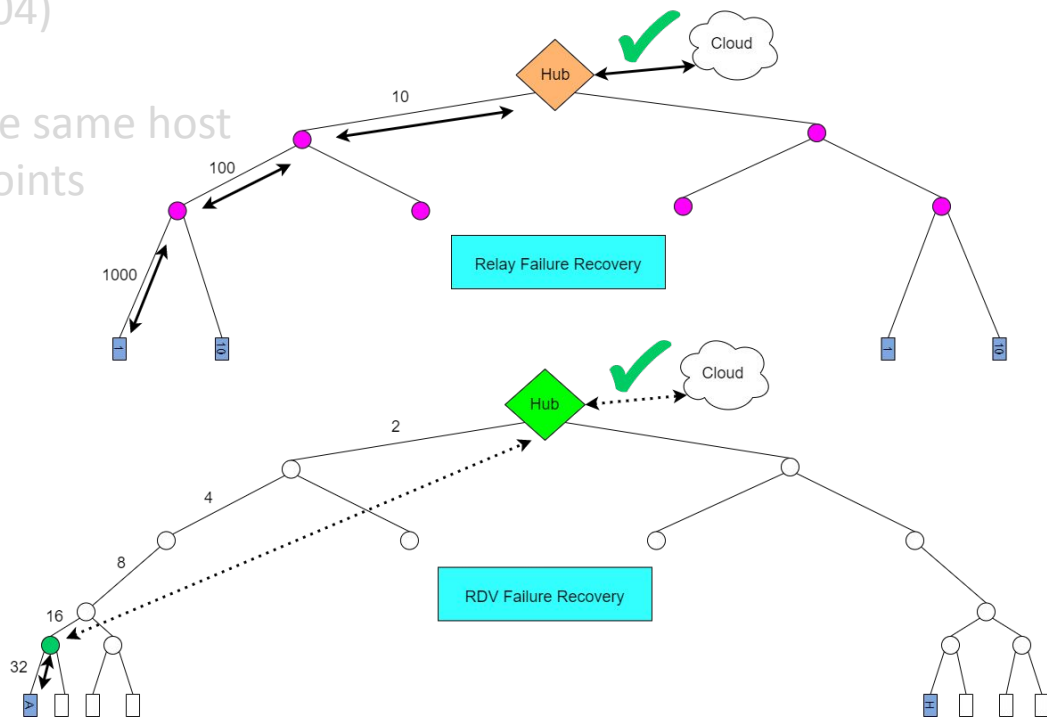
# Experiments

## • Setup

- Docker Linux container (Ubuntu 20.04)
- RDV node coverage: 10%
- Run components as processes on the same host
- Scalability analysis up to 667K endpoints

## • Metrics


- Intermediate nodes
- Total physical edges
- **Control/data messages**



# Analysis: Minimizing Intermediate Nodes

## Minimizing intermediate nodes


- **Experiment:**
  - To estimate the number of intermediate nodes required to get messages to the hub
- **Observation(s):**
  - RDV method has **10-1000x fewer intermediate nodes** depending on targeted coverage
  - RDV selection algorithm is **sensitive** to coverage and costs



# Analysis: Detecting and Adapting to Failures

## Time to acquiesce/adapt to failures


- **Experiment:**
  - To understand the number of messages required to update all endpoints of single point of failure (i.e., the hub)
- **Observation(s):**
  - RDV failure messages **scale gracefully** with number of endpoints



# Analysis: Recovering from Failures

## Model download in failure recovery

- **Experiment:**
  - To understand the number of control/data messages required to pull a new update during failure recovery
- **Observation(s):**
  - RDV method **scales constantly** with number of endpoints

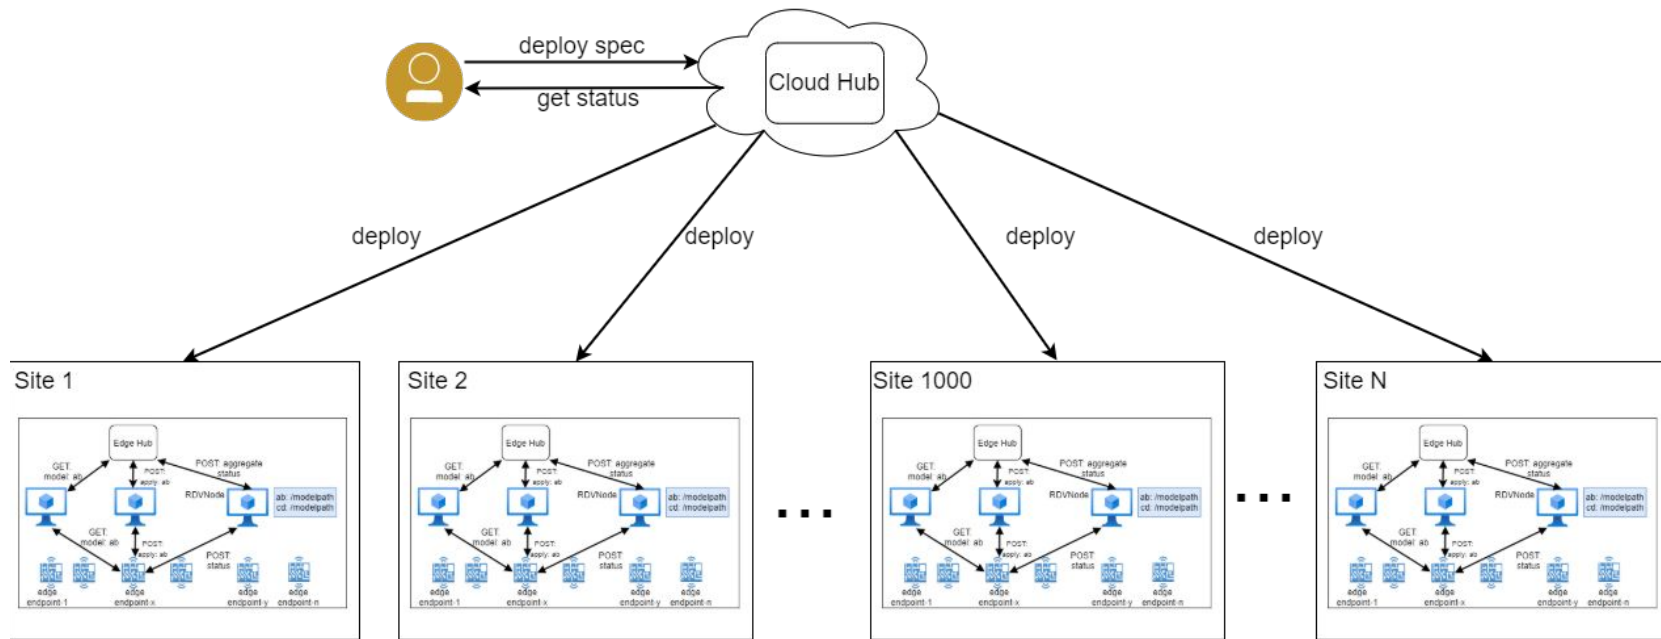




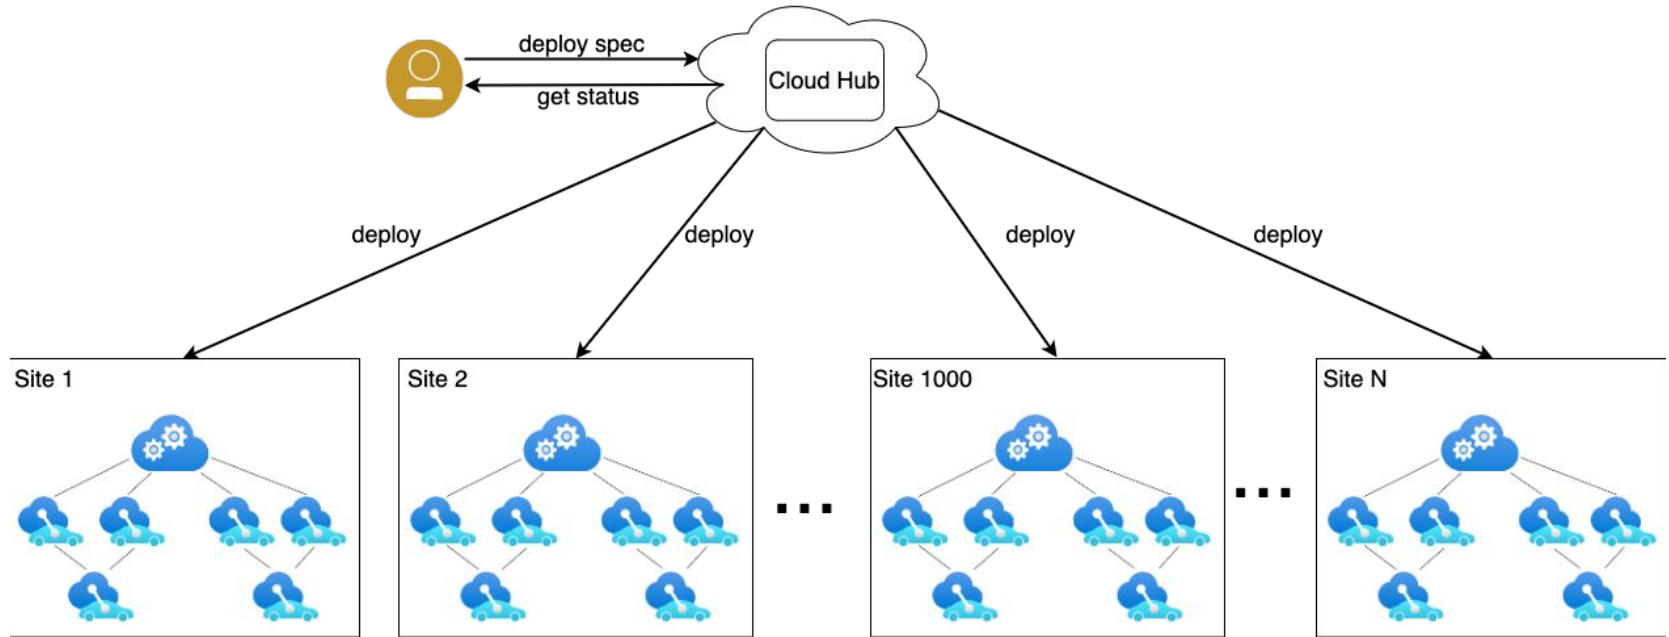
# Next Steps

- Hardware implementation
  - Scalability
  - Bandwidth consumption
  - Data transfer
  - RTT optimization
- Novel applications
  - Digital agriculture


# Conclusion



# Conclusion



# Conclusion



- 1-3 orders of magnitude fewer intermediate nodes
- Scalable infrastructure bootstrapping
- Adjustable network resilience
- Efficient resource usage

# Thank You



Gloire Rubambiza

[gloire@cs.cornell.edu](mailto:gloire@cs.cornell.edu)

<https://rubambiza.github.io>